

```

!-----
! Skeleton 2D Electrostatic GPU-MPI PIC code
! written by Viktor K. Decyk, UCLA
    program gpufppic2
    use fgpulib2
    use fgpuppush2
    use fgpupfft2
    use ppush2_h
    use pplib2          ! use with pplib2.f90
!   use pplib2_h        ! use with pplib2.f
    use fgpplib2
    use dtimer_c
    implicit none
    integer, parameter :: indx = 9, indy = 9
    integer, parameter :: npx = 3072, npy = 3072
    integer, parameter :: ndim = 2
    real, parameter :: tend = 10.0, dt = 0.1, qme = -1.0
    real, parameter :: vtx = 1.0, vty = 1.0, vx0 = 0.0, vy0 = 0.0
    real :: ax = .912871, ay = .912871
! idimp = dimension of phase space = 4
    integer :: idimp = 4, ipbc = 1
! idps = number of partition boundaries
    integer :: idps = 2
    real :: wke = 0.0, we = 0.0, wt = 0.0
! sorting tiles
    integer :: mx = 16, my = 16
! fraction of extra particles needed for particle management
    real :: xtras = 0.2
! declare scalars for standard code
    integer :: nx, ny, nxh, nyh, nxh1, nxe, nye, nxeh
    integer :: nxyh, nxhy, mx1, ntime, nloop, isign, ierr
    real :: qbme, affp
    real, dimension(1) :: ssum
    double precision :: np
!
! declare scalars for MPI code
    integer :: nvp, idproc, kstrt, npmax, kxp, kyp, nypmx, nypmn
    integer :: nyp, noff, npp, nps
    integer :: myp1, mxyp1, kxpp, kypp
!
! declare scalars for GPU code
    integer :: nblock = 128
! nscache = (0,1,2) = (no,small,big) cache size
    integer :: nscache = 1
    integer :: mmcc, nppmx, nppmx0, nbmaxp, ntmaxp, npbm
    integer :: nxhd, kxpd, idev, ndev
    integer, dimension(1) :: irc
!
! declare arrays for standard code
    real, dimension(:,:), pointer :: part => null()
    complex, dimension(:,:), pointer :: ffct => null()
    integer, dimension(:), pointer :: mixup => null()
    complex, dimension(:), pointer :: sct => null()
    real, dimension(4) :: wtot, work

```

```

!
! declare arrays for MPI code
    real, dimension(:), pointer :: edges => null()
    integer, dimension(:), pointer :: locl => null()
    complex, pinned, dimension(:), allocatable :: scs, scr
    real, pinned, dimension(:), allocatable :: sbuf1, sbufr
    real, pinned, dimension(:), allocatable :: rbuf1, rbufr
!
! declare arrays for GPU code
    real, device, dimension(:,:), allocatable :: g_ge
    real, device, dimension(:,:,:), allocatable :: g_fxye
    complex, device, dimension(:,:), allocatable :: g_q, g_qt
    complex, device, dimension(:,:,:), allocatable :: g_fxy, g_fxyt
    complex, device, dimension(:,:), allocatable :: g_ffct
    integer, device, dimension(:), allocatable :: g_mixup
    complex, device, dimension(:), allocatable :: g_sct
    real, device, dimension(:), allocatable :: g_wke, g_we
    real, device, dimension(:,:,:), allocatable :: g_ppart, g_ppbuff
    integer, device, dimension(:), allocatable :: g_kpic
    real, device, dimension(:), allocatable :: g_sbuf1, g_sbufr
    integer, device, dimension(:,:), allocatable :: g_ncl
    integer, device, dimension(:,:,:), allocatable :: g_ihole
    integer, device, dimension(:,:), allocatable :: g_nc11, g_nclr
    complex, device, dimension(:,:), allocatable :: g_bsm, g_brm
    complex, device, dimension(:), allocatable :: g_scs
    real, device, dimension(:), allocatable :: g_sum
    integer, device, dimension(:), allocatable :: g_irc
    complex, dimension(:,:), pointer :: qt => null()
    complex, dimension(:,:,:), pointer :: fxyt => null()
    real, dimension(:,:,:), pointer :: ppart => null()
    integer, dimension(:), pointer :: kp1c => null()
    integer, dimension(:,:), pointer :: nc11 => null(), nclr => null()
    integer, dimension(:,:), pointer :: mc11 => null(), mclr => null()
    complex, pinned, dimension(:,:), allocatable :: bsm, brm
!
! declare and initialize timing data
    real :: time
    type (timeval) :: itime
    double precision :: dtime
    real :: tpush = 0.0, tdpost = 0.0, tsort = 0.0
    real :: tmov = 0.0, tfield = 0.0, tguard = 0.0
    real, dimension(2) :: tmsort = 0.0
    real, dimension(2) :: tfft = 0.0
!
! initialize scalars for standard code
    np = dble(np1)*dble(np2)
    nx = 2**indx; ny = 2**indy; nxh = nx/2; nyh = ny/2
    nxh1 = nxh + 1
    nxe = nx + 2; nye = ny + 2; nxeh = nxe/2
    nxyh = max(nx,ny)/2; nxhy = max(nxh,ny)
    mx1 = (nx - 1)/mx + 1
    nloop = tend/dt + .0001; ntime = 0
    qbme = qme
    affp = dble(nx)*dble(ny)/np

```

```

! set size for FFT arrays
    nxhd = nxh1
!
! nvp = number of distributed memory nodes
! initialize for distributed memory parallel processing
    call PPINIT2(idproc,nvp)
    kstrt = idproc + 1
! check if too many processors
    if (nvp > ny) then
        if (kstrt==1) then
            write (*,*) 'Too many processors requested: ny, nvp=', ny, nvp
        endif
        go to 3000
    endif
!
! initialize data for MPI code
    allocate(edges(idps))
! calculate partition variables: edges, nyp, noff, nypmx
! edges(1:2) = lower:upper boundary of particle partition
! nyp = number of primary (complete) gridpoints in particle partition
! noff = lowermost global gridpoint in particle partition
! nypmx = maximum size of particle partition, including guard cells
! nypmn = minimum value of nyp
    call PDICOMP2L(edges,nyp,noff,nypmx,nypmn,ny,kstrt,nvp,idps)
    if (nypmn < 1) then
        if (kstrt==1) then
            write (*,*) 'combination not supported nvp, ny =',nvp,ny
        endif
        go to 3000
    endif
! initialize additional scalars for MPI code
! kxp = number of complex grids in each field partition in x direction
! kxpd = number of complex grids in fft partition in x direction
    kxp = (nxh - 1)/nvp + 1
! set size for FFT arrays
    kxpd = (nxh1 - 1)/nvp + 1
! kyp = number of complex grids in each field partition in y direction
    kyp = (ny - 1)/nvp + 1
! npmax = maximum number of electrons in each partition
    npmax = (np/nvp)*1.25
! myp1 = number of tiles in y direction
    myp1 = (nyp - 1)/my + 1; mxyp1 = mx1*myy1
! kxpp/kypp = actual size of GPU field partition
    kxpp = min(kxpd,max(0,nxhd-kxpd*idproc))
    kypp = min(kyp,max(0,ny-kyp*idproc))
!
! allocate and initialize data for standard code
    allocate(part(idimp,npmax))
    allocate(ffct(nyh,kxpd))
    allocate(mixup(nxhy),sct(nxyh))
    allocate(kpic(mxyp1))
    allocate(qt(ny,kxpd),fxyt(ny,ndim,kxpd))
!
! allocate and initialize data for MPI code

```

```

        allocate(locl(nvp))
!
! set up GPU
    irc = 0
! get unique GPU device ids
    call PPFNDGRP(locl,kstrt,nvp,idev,ndev)
    if (idev < 0) then
        write (*,*) kstrt,'GPU device id error!'
        call PPABORT()
        stop
    endif
    call fgpu_setgbsize(nblock)
    call init_cuf(idev,irc(1))
    if (irc(1) /= 0) then
        write (*,*) kstrt,'CUDA initialization error!'
        call PPABORT()
        stop
    endif
! obtain compute capability
    mmcc = fgetmmcc()
    if (mmcc < 20) then
        write (*,*) kstrt, 'compute capability 2.x or higher required'
        call PPABORT()
        stop
    endif
! set cache size
    call fgpu_set_cache_size(nscache)
! create asynchronous streams
    call fgpu_initstream(1)
    call fgpu_initstream(2)
    call fgpu_initstream(3)
! allocate and initialize data for GPU code
    allocate(g_ge(nxe,nypmx),g_fxye(ndim,nxe,nypmx))
    allocate(g_ffct(nyh,kxpd),g_mixup(nxhy),g_sct(nxyh))
    allocate(g_q(nxhd,kyp),g_qt(ny,kxpd))
    allocate(g_fxy(nxhd,ndim,kyp),g_fxyt(ny,ndim,kxpd))
    allocate(g_wke(mxyp1),g_we(kxpd))
    allocate(g_sum(1))
!
! prepare fft tables
    call WPPFFT2RINIT(mixup,sct,indx,indy,nxhy,nxyh)
! prepare NVIDIA ffts
    call fgpupfft2rrcuinit(nx,kypp,ndim)
    call fgpupfft2cuinit(kxpp,ny,ndim)
! calculate form factors
    isign = 0
    call PPOIS22T(qt,fxyt,isign,ffct,ax,ay,affp,we,nx,ny,kstrt,ny,kxpd&
        &,nyh)
! copy in solver arrays to GPU
    g_mixup = mixup
    g_sct = sct
    g_ffct = ffct
! initialize electrons
    nps = 1

```

```

    npp = 0
    call PDISTR2(part,edges,npp,nps,vtx,vty,vx0,vy0,npx,nty,nx,ny,      &
&idimp,npmax,idps,ipbc,ierr)
! check for particle initialization error
    if (ierr /= 0) then
        if (kstrt==1) then
            write (*,*) 'particle initialization error: ierr=', ierr
        endif
        go to 3000
    endif

!
! find number of particles in each of mx, my tiles: updates kplic, nppmx
    call PPDBLK2L(part,kpic,npp,noff,nppmx,idimp,npmax,mx,my,mx1,      &
&mxypl,irc(1))
    if (irc(1) /= 0) then
        write (*,*) kstrt, 'PPDBLK2L error, irc=', irc(1)
        call PPABORT()
        stop
    endif

! allocate vector particle data
    nppmx0 = (1.0 + xtras)*nppmx
    ntmaxp = xtras*nppmx
    npbm = xtras*nppmx
    nbmaxp = 0.25*mx1*npbm

! align data to warp size
    nppmx0 = 32*((nppmx0 - 1)/32 + 1)
    ntmaxp = 32*(ntmaxp/32 + 1)
    npbm = 32*((npbm - 1)/32 + 1)
    nbmaxp = 32*((nbmaxp - 1)/32 + 1)
    allocate(g_ppart(nppmx0,idimp,mxypl),g_ppbuff(npbm,idimp,mxypl))
    allocate(g_kpic(mxypl))
    allocate(g_sbuf1(nbmaxp*idimp),g_sbufr(nbmaxp*idimp))
    allocate(g_ncl(8,mxypl),g_ihole(2,ntmaxp+1,mxypl))
    allocate(g_nc11(3,mx1),g_nclr(3,mx1))
    allocate(g_bsm(kxpd*ndim*kyp,nvp),g_brm(kxpd*ndim*kyp,nvp))
    allocate(g_scs(nxeh*ndim))
    allocate(g_irc(1))
    allocate(ppart(nppmx0,idimp,mxypl))
    allocate(nc11(3,mx1),nclr(3,mx1),mc11(3,mx1),mclr(3,mx1))

!
! allocate host page-locked memory for GPU-MPI buffers
    allocate(scs(nxeh*ndim),scr(nxeh*ndim))
    allocate(sbuf1(idimp*nbmaxp),sbufr(idimp*nbmaxp))
    allocate(rbuf1(idimp*nbmaxp),rbufr(idimp*nbmaxp))
    allocate(bsm(kxpd*ndim*kyp,nvp),brm(kxpd*ndim*kyp,nvp))

!
! copy ordered particle data for GPU code: updates ppart, kplic
    call PPPMOVIN2LT(part,ppart,kpic,npp,noff,nppmx0,idimp,npmax,mx,my&
&mx1,mxypl,irc(1))
    if (irc(1) /= 0) then
        write (*,*) kstrt, 'PPPMOVIN2LT overflow error, irc=', irc(1)
        call PPABORT()
        stop
    endif
endif

```

```

! sanity check
    call PPPCHECK2LT(ppart,kpic,noff,nyp,idimp,nppmx0,nx,mx,my,mx1,    &
    &mypl,irc(1))
    if (irc(1) /= 0) then
        write (*,*) kstrt, 'PPPCHECK2LT error: irc=', irc(1)
        call PPABORT()
        stop
    endif
! copy to GPU
    g_irc = irc
    g_ppart = ppart
    g_kpic = kpic
    call fgpu_zfmem(g_we,kxpd)
!
! * * * start main iteration loop * * *
!
500 if (nloop <= ntime) go to 2000
!     if (kstrt==1) write (*,*) 'ntime = ', ntime
!
! deposit charge with GPU code: updates g_ge
    call dtimer(dtime,itime,-1)
    call fgpu_zfmem(g_ge,nxe*nypmx)
    call fgpu2ppgppost2l(g_ppart,g_ge,g_kpic,noff,qme,idimp,nppmx0,mx,&
    &my,nxe,nypmx,mx1,mxyp1)
    call dtimer(dtime,itime,1)
    time = real(dtime)
    tdpot = tdpot + time
!
! add and copy guard cells with GPU code: updates g_q
    call dtimer(dtime,itime,-1)
    call GPPCAGUARD2L(g_q,g_ge,g_scs,scs,scr,nx,nyp,kstrt,nvp,nxe,    &
    &nypmx,nxhd,kyp)
    call dtimer(dtime,itime,1)
    time = real(dtime)
    tguard = tguard + time
!
! transform charge to fourier space with GPU code: updates g_q, g_qt,
! as well as various buffers
    isign = -1
    call WAPPPFT2RCS(g_q,g_qt,g_bsm,g_brm,bsm,brm,isign,g_mixup,g_sct,&
    &tfft,indx,indy,kstrt,nvp,kxpd,kyp,nxhd,ny,kyp,nxhy,nxyh)
! NVIDIA fft
!     call GPUPPFT2RRCU(g_q,g_qt,g_bsm,g_brm,bsm,brm,isign,tfft,indx,    &
!     &indy,kstrt,nvp,kxpd,kyp,nxhd,ny,kyp)
!
! calculate force/charge in fourier space with GPU code:
! updates g_fxyt, g_we
    call dtimer(dtime,itime,-1)
    call fgpuppois22t(g_qt,g_fxyt,g_ffct,g_we,nx,ny,kstrt,ny,kxpd,nyh)
    call dtimer(dtime,itime,1)
    time = real(dtime)
    tfield = tfield + time
!
! transform force to real space with GPU code: updates g_fxy, g_fxyt,

```

```

! as well as various buffers
    isign = 1
    call WAPPPFFT2RCSN(g_fxy,g_fxyt,g_bsm,g_brm,bsm,brm,isign,g_mixup, &
        &g_sct,tfft,indx,indy,kstrt,nvp,ndim,kxpd,kyp,nxhd,ny,kyp,nxhy,nxyh&
        &)
! NVIDIA fft
!     call GPUPPFFT2RRCUN(g_fxy,g_fxyt,g_bsm,g_brm,bsm,brm,isign,tfft,  &
!     &indx,indy,kstrt,nvp,ndim,kxpd,kyp,nxhd,ny,kyp)
!
! copy guard cells with GPU code: updates g_fxye
    call dtimer(dtime,itime,-1)
    call GPPCCGUARD2L(g_fxy,g_fxye,g_scs,scs,scr,nx,nyp,kstrt,nvp,ndim&
        &,nxe,nypmx,nxhd,kyp)
    call dtimer(dtime,itime,1)
    time = real(dtime)
    tguard = tguard + time
!
! push particles with GPU code: updates g_ppart, g_wke
    call dtimer(dtime,itime,-1)
    call fgpuppjpgpush2l(g_ppart,g_fxye,g_kpic,noff,nyp,qbme,dt,g_wke, &
        &nx,ny,mx,my,idimp,nppmx0,nxe,nypmx,mx1,mxyp1,ipbc)
    call dtimer(dtime,itime,1)
    time = real(dtime)
    tpush = tpush + time
!
! reorder particles by tile with GPU code:
! updates g_ppart, g_ppbuff, g_kpic, g_ncl, g_ihole, and g_irc,
! as well as various buffers
    call GPPORDER2L(g_ppart,g_ppbuff,g_sbuf1,g_sbuf2,g_kpic,g_ncl,      &
        &g_ihole,g_ncl1,g_ncl2,sbuf1,sbuf2,rbuf1,rbuf2,ncl1,ncl2,mcl1,mcl2,&
        &tmsort,noff,nyp,kstrt,nvp,idimp,nppmx0,nx,ny,mx,my,mx1,myp1,npbmx,&
        &ntmaxp,nbmaxp,g_irc)
    tsort = tmsort(1)
    tmov = tmsort(2)
!
! sanity check
    irc = g_irc(1)
    if (irc(1) /= 0) then
        write (*,*) kstrt, 'GPPORDER2L error: irc=', irc(1)
        call PPABORT()
        stop
    endif
!
! energy diagnostic
    if (ntime==0) then
        call fgpu_zfmem(g_sum,1)
        call fgpsum2(g_we,g_sum,kxpd)
        we = g_sum(1)
        call fgpu_zfmem(g_sum,1)
        call fgpsum2(g_wke,g_sum,mxyp1)
        wke = g_sum(1)
        wtot(1) = we
        wtot(2) = wke
        wtot(3) = 0.0
    
```

```

    wtot(4) = we + wke
    call PPSUM(wtot,work,4)
    we = wtot(1)
    wke = wtot(2)
    if (kstrt==1) then
        write (*,*) 'Initial Field, Kinetic and Total Energies:'
        write (*,'(3e14.7)') we, wke, wke + we
    endif
endif
ntime = ntime + 1
go to 500
2000 continue
!
! * * * end main iteration loop * * *
!
! energy diagnostic
    call fgpu_zfmem(g_sum,1)
    call fgpusum2(g_we,g_sum,kxpd)
    we = g_sum(1)
    call fgpu_zfmem(g_sum,1)
    call fgpusum2(g_wke,g_sum,mxyp1)
    wke = g_sum(1)
    wtot(1) = we
    wtot(2) = wke
    wtot(3) = 0.0
    wtot(4) = we + wke
    call PPSUM(wtot,work,4)
    we = wtot(1)
    wke = wtot(2)
!
    if (kstrt==1) then
        write (*,*) 'ntime = ', ntime
        write (*,*) 'MPI nodes nvp = ', nvp, ', GPUs per host = ', ndev
        write (*,*) 'Final Field, Kinetic and Total Energies:'
        write (*,'(3e14.7)') we, wke, wke + we
        write (*,*)
!
        write (*,*) 'deposit time = ', tdpst
        write (*,*) 'guard time = ', tguard
        write (*,*) 'solver time = ', tfield
        write (*,*) 'fft times = ', sum(tfft), tfft
        write (*,*) 'push time = ', tpush
        write (*,*) 'move time = ', tmov
        write (*,*) 'sort time = ', tsort
        tfield = tfield + tguard + sum(tfft)
        write (*,*) 'total solver time = ', tfield
        time = tdpst + tpush + tsort + tmov
        write (*,*) 'total particle time = ', time
        wt = time + tfield
        write (*,*) 'total time = ', wt
        write (*,*)
!
        wt = 1.0e+09/(real(nloop)*real(np))
        write (*,*) 'Push Time (nsec) = ', tpush*wt

```



```

        write (*,*) 'Deposit Time (nsec) = ', tdpost*wt
        write (*,*) 'Sort Time (nsec) = ', tsort*wt
        write (*,*) 'Move Time (nsec) = ', tmov*wt
        write (*,*) 'Total Particle Time (nsec) = ', time*wt
    endif
!
! close down NVIDIA fft
    call fgpupfft2cudel()
    call fgpupfft2rrcudel()
! deallocate memory on GPU
    deallocate(g_irc,g_scs,g_brm,g_bsm,g_nclr,g_ncll,g_ihole,g_ncl)
    deallocate(g_sbufr,g_sbufl,g_kpic,g_ppbuff,g_ppart)
    deallocate(g_sum,g_we,g_wke)
    deallocate(g_fxyt,g_fxy,g_qt,g_q,g_sct,g_mixup,g_ffct,g_fxye,g_ge)
! deallocate host page-locked memory
    deallocate(scs,scr,sbufl,sbufr,rbufl,rbufr,bsm,brm)
3000 continue
!
! delete asynchronous streams
    call fgpu_delstream(3)
    call fgpu_delstream(2)
    call fgpu_delstream(1)
! close down GPU
    call end_cuf()
! close down MPI
    call PPEXIT()
!
    stop
end program

```