

```

/*-----*/
__global__ void gpuppsh2l(float ppart[], float fxy[], int kplic[],
                        float qbm, float dt, float *ek, int idimp,
                        int nppmx, int nx, int ny, int mx, int my,
                        int nxv, int nyv, int mx1, int mxy1,
                        int ipbc) {
/* for 2d code, this subroutine updates particle co-ordinates and
   velocities using leap-frog scheme in time and first-order linear
   interpolation in space, with various boundary conditions.
   threaded version using guard cells
   data read in tiles
   particles stored segmented array
   44 flops/particle, 12 loads, 4 stores
   input: all, output: ppart, ek
   equations used are:
   vx(t+dt/2) = vx(t-dt/2) + (q/m)*fx(x(t),y(t))*dt,
   vy(t+dt/2) = vy(t-dt/2) + (q/m)*fy(x(t),y(t))*dt,
   where q/m is charge/mass, and
   x(t+dt) = x(t) + vx(t+dt/2)*dt, y(t+dt) = y(t) + vy(t+dt/2)*dt
   fx(x(t),y(t)) and fy(x(t),y(t)) are approximated by interpolation from
   the nearest grid points:
   fx(x,y) = (1-dy)*((1-dx)*fx(n,m)+dx*fx(n+1,m)) + dy*((1-dx)*fx(n,m+1)
               + dx*fx(n+1,m+1))
   fy(x,y) = (1-dy)*((1-dx)*fy(n,m)+dx*fy(n+1,m)) + dy*((1-dx)*fy(n,m+1)
               + dx*fy(n+1,m+1))
   where n,m = leftmost grid points and dx = x-n, dy = y-m
   ppart[m][0][n] = position x of particle n in tile m
   ppart[m][1][n] = position y of particle n in tile m
   ppart[m][2][n] = velocity vx of particle n in tile m
   ppart[m][3][n] = velocity vy of particle n in tile m
   fxy[k][j][0] = x component of force/charge at grid (j,k)
   fxy[k][j][1] = y component of force/charge at grid (j,k)
   that is, convolution of electric field over particle shape
   kplic = number of particles per tile
   qbm = particle charge/mass
   dt = time interval between successive calculations
   kinetic energy/mass at time t is also calculated, using
   ek = .125*sum((vx(t+dt/2)+vx(t-dt/2))*2+(vy(t+dt/2)+vy(t-dt/2))*2)
   idimp = size of phase space = 4
   nppmx = maximum number of particles in tile
   nx/ny = system length in x/y direction
   mx/my = number of grids in sorting cell in x/y
   nxv = first dimension of field arrays, must be >= nx+1
   nyv = second dimension of field arrays, must be >= ny+1
   mx1 = (system length in x direction - 1)/mx + 1
   mxy1 = mx1*my1, where my1 = (system length in y direction - 1)/my + 1
   ipbc = particle boundary condition = (0,1,2,3) =
   (none,2d periodic,2d reflecting,mixed reflecting/periodic)
local data
int noff, moff, npoff, npp, mxv;
int i, j, k, ii, nn, mm;
float qtm, edgelx, edgely, edgerx, edgery, dxp, dyp, amx, amy;
float x, y, dx, dy, vx, vy;
/* The sizes of the shared memory arrays are as follows: */

```

```

/* float sfxy[2*(mx+1)*(my+1)], swke[blockDim.x];          */
/* to conserve memory, swke overlaps with sfxy             */
/* and the name sfxy is used instead of swke                */
extern __shared__ float sfxy[];
double sum1;
qtm = qbm*dt;
sum1 = 0.0;
/* set boundary values */
edgelx = 0.0f;
edgely = 0.0f;
edgerx = (float) nx;
edgery = (float) ny;
if (ipbc==2) {
    edgelx = 1.0f;
    edgely = 1.0f;
    edgerx = (float) (nx-1);
    edgery = (float) (ny-1);
}
else if (ipbc==3) {
    edgelx = 1.0f;
    edgerx = (float) (nx-1);
}
mxv = mx + 1;
/* k = tile number */
k = blockIdx.x + gridDim.x*blockIdx.y;
/* loop over tiles */
if (k < mxy1) {
    noff = k/mx1;
    moff = my*noff;
    noff = mx*(k - mx1*noff);
    npp = kplic[k];
    npoff = idimp*nppmx*k;
/* load local fields from global array */
    nn = (mx < nx-noff ? mx : nx-noff) + 1;
    mm = (my < ny-moff ? my : ny-moff) + 1;
    ii = threadIdx.x;
    while (ii < mxv*(my+1)) {
        j = ii/mxv;
        i = ii - mxv*j;
        if ((i < nn) && (j < mm)) {
            sfxy[2*ii] = fxy[2*(i+noff+nxv*(j+moff))];
            sfxy[1+2*ii] = fxy[1+2*(i+noff+nxv*(j+moff))];
        }
        ii += blockDim.x;
    }
/* synchronize threads */
    __syncthreads();
/* loop over particles in tile */
    j = threadIdx.x;
    while (j < npp) {
/* find interpolation weights */
        x = ppart[j+npoff];
        nn = x;
        y = ppart[j+npoff+nppmx];
    }
}

```

```

    mm = y;
    dxp = x - (float) nn;
    dyp = y - (float) mm;
    nn = 2*(nn - noff) + 2*mxv*(mm - moff);
    amx = 1.0f - dxp;
    amy = 1.0f - dyp;
/* find acceleration */
    dx = amx*sfxxy[nn];
    dy = amx*sfxxy[1+nn];
    dx = amy*(dxp*sfxxy[2+nn] + dx);
    dy = amy*(dyp*sfxxy[3+nn] + dy);
    nn += 2*mxv;
    vx = amx*sfxxy[nn];
    vy = amx*sfxxy[1+nn];
    dx += dyp*(dxp*sfxxy[2+nn] + vx);
    dy += dyp*(dyp*sfxxy[3+nn] + vy);
/* new velocity */
    vx = ppart[j+npoff+nppmx*2];
    vy = ppart[j+npoff+nppmx*3];
    dx = vx + qtm*dx;
    dy = vy + qtm*dy;
/* average kinetic energy */
    vx += dx;
    vy += dy;
    sum1 += (double) (vx*vx + vy*vy);
    ppart[j+npoff+nppmx*2] = dx;
    ppart[j+npoff+nppmx*3] = dy;
/* new position */
    dx = x + dx*dt;
    dy = y + dy*dt;
/* reflecting boundary conditions */
    if (ipbc==2) {
        if ((dx < edgelx) || (dx >= edgerx)) {
            dx = ppart[j+npoff];
            ppart[j+npoff+nppmx*2] = -ppart[j+npoff+nppmx*2];
        }
        if ((dy < edgely) || (dy >= edgery)) {
            dy = ppart[j+npoff+nppmx];
            ppart[j+npoff+nppmx*3] = -ppart[j+npoff+nppmx*3];
        }
    }
/* mixed reflecting/periodic boundary conditions */
    else if (ipbc==3) {
        if ((dx < edgelx) || (dx >= edgerx)) {
            dx = ppart[j+npoff];
            ppart[j+npoff+nppmx*2] = -ppart[j+npoff+nppmx*2];
        }
    }
/* set new position */
    ppart[j+npoff] = dx;
    ppart[j+npoff+nppmx] = dy;
    j += blockDim.x;
}
__syncthreads();

```

```

/* add kinetic energies in tile */
    sfxxy[threadIdx.x] = (float) sum1;
/* synchronize threads */
    __syncthreads();
    lsum2(sfxxy,blockDim.x);
/* normalize kinetic energy of tile */
    if (threadIdx.x==0) {
        ek[k] = 0.125f*sfxxy[0];
    }
}
return;
}

```

```

/*-----*/
__global__ void gpu2ppost2l(float ppart[], float q[], int kplic[],
                           float qm, int nppmx, int idimp, int mx,
                           int my, int nxv, int nyv, int mx1,
                           int mxy1) {
/* for 2d code, this subroutine calculates particle charge density
using first-order linear interpolation, periodic boundaries
threaded version using guard cells
data deposited in tiles
particles stored segmented array
17 flops/particle, 6 loads, 4 stores
input: all, output: q
charge density is approximated by values at the nearest grid points
q(n,m)=qm*(1.-dx)*(1.-dy)
q(n+1,m)=qm*dx*(1.-dy)
q(n,m+1)=qm*(1.-dx)*dy
q(n+1,m+1)=qm*dx*dy
where n,m = leftmost grid points and dx = x-n, dy = y-m
ppart[m][0][n] = position x of particle n in tile m
ppart[m][1][n] = position y of particle n in tile m
q[k][j] = charge density at grid point j,k
kplic = number of particles per tile
qm = charge on particle, in units of e
nppmx = maximum number of particles in tile
idimp = size of phase space = 4
mx/my = number of grids in sorting cell in x/y
nxv = first dimension of charge array, must be >= nx+1
nyv = second dimension of charge array, must be >= ny+1
mx1 = (system length in x direction - 1)/mx + 1
mxy1 = mx1*my1, where my1 = (system length in y direction - 1)/my + 1
local data
int noff, moff, npoff, npp, mxv;
int i, j, k, ii, nn, mm, np, mp;
float dxp, dyp, amx, amy;
extern __shared__ float sq[];
mxv = mx + 1;
k = blockIdx.x + gridDim.x*blockIdx.y;
/* loop over tiles */
if (k < mxy1) {
    noff = k/mx1;
    moff = my*noff;
    noff = mx*(k - mx1*noff);
    npp = kplic[k];
    npoff = idimp*nppmx*k;
/* zero out local accumulator */
    i = threadIdx.x;
    while (i < mxv*(my+1)) {
        sq[i] = 0.0f;
        i += blockDim.x;
    }
/* synchronize threads */
    __syncthreads();
/* loop over particles in tile */

```

```

        j = threadIdx.x;
        while (j < npp) {
/* find interpolation weights */
            dxp = ppart[j+npoff];
            nn = dxp;
            dyp = ppart[j+npoff+nppmx];
            mm = dyp;
            dxp = qm*(dxp - (float) nn);
            dyp = dyp - (float) mm;
            nn = nn - noff;
            mm = mxv*(mm - moff);
            amx = qm - dxp;
            mp = mm + mxv;
            amy = 1.0f - dyp;
            np = nn + 1;
/* deposit charge within tile to local accumulator */
/* original deposit charge, has data hazard on GPU */
/*      sq[np+mp] += dxp*dyp; */
/*      sq[nn+mp] += amx*dyp; */
/*      sq[np+mm] += dxp*amy; */
/*      sq[nn+mm] += amx*amy; */
/* for devices with compute capability 2.x */
            atomicAdd(&sq[np+mp],dxp*dyp);
            atomicAdd(&sq[nn+mp],amx*dyp);
            atomicAdd(&sq[np+mm],dxp*amy);
            atomicAdd(&sq[nn+mm],amx*amy);
            j += blockDim.x;
        }
/* synchronize threads */
        __syncthreads();
/* deposit charge to global array */
        nn = mxv < nxv-noff ? mxv : nxv-noff;
        mm = my+1 < nyv-moff ? my+1 : nyv-moff;
        ii = threadIdx.x;
        while (ii < mxv*(my+1)) {
            j = ii/mxv;
            i = ii - mxv*j;
            if ((i < nn) && (j < mm)) {
/* original deposit charge, has data hazard on GPU */
/*      q[i+noff+nxv*(j+moff)] += sq[ii]; */
/* for devices with compute capability 1.x */
/*      gatomicAdd(&q[i+noff+nxv*(j+moff)],sq[ii]); */
/* for devices with compute capability 2.x */
            atomicAdd(&q[i+noff+nxv*(j+moff)],sq[ii]);
            }
            ii += blockDim.x;
        }
    }
    return;
}

```